

Independent Task Scheduling

Embedded OS Implementation

Prof. Ya-Shu Chen


National Taiwan University
of Science and Technology

Motivation


- Take yourself as an example
 - Naturally you have a number of things to do with time pressure
 - Project deadlines, meeting time, class time, and deadlines for bills
 - Some of them regularly recur but some don't
 - To eat meal on 12:30 everyday
 - Go to the movies on 8:00pm

Motivation

- You schedule yourself to meet deadlines
 - Course A: one homework is announced every 9 days, each costs you 4.5 days to do
 - Course B: one homework is announced every 6 days, each costs you 3 days to do
- You miss deadlines of one course if your policy **favors either one course**




Favor course B



Favor course A


Motivation

- Schedule to meet deadlines (cont'd.)
 - Course A: (4.5, 9)
 - Course B: (3, 6)
- All deadlines are met if you whatever has the **closest deadline**




Motivation

- You schedule yourself to survive overloads
 - $(1,2)$, $(1,3)$, $(1,1.5)$



Favoring $(1,1.5) \rightarrow (1,3) \rightarrow (1,2)$
2 lecturers are happy, 1 will flunk you though...




Do whatever has the closest deadline.
You are in deep shit!

Cyclic-Executive

Cyclic Executive

pros: 方便degug
cons: 難以修改

- The system repeatedly exercises a static schedule
 - A table-driven approach
- Many existing systems still take this approach
 - Easy to debug and easy to visualize
 - Highly deterministic
 - Hard to program, to modify, and to upgrade
 - A program should be divided into many pieces (like an FSM)



Cyclic Executive

- The table emulates an infinite loop of routines
 - However, a single independent loop is not enough to many complicated systems
 - Multiple concurrent loops should be considered
- How large should the table be when there are multiple loops?



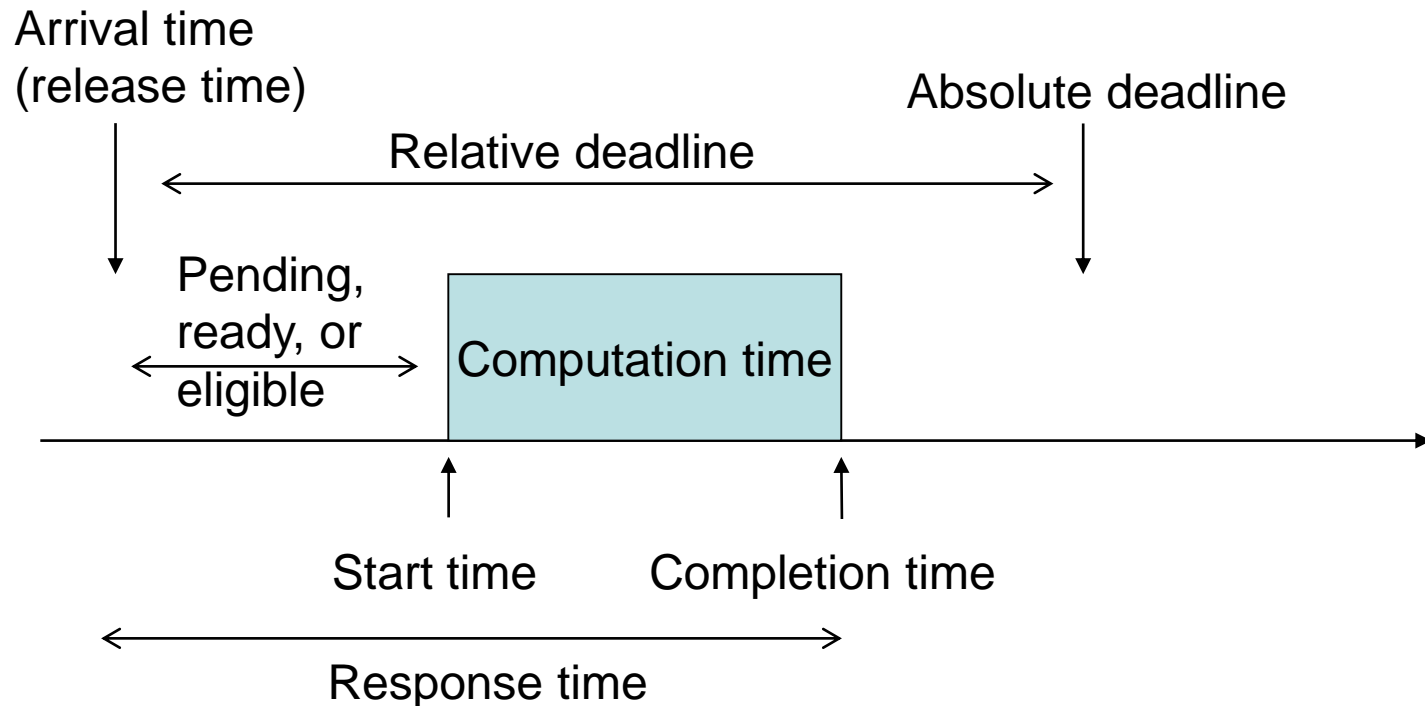
Cyclic Executive

- **Definition:** Let the **hyper-period** of a collection of loops be a time interval which's length is the least-common-multiplier of the loops' lengths
 - Let the length of the hyper-period be abbreviated as “h”
- **Theorem:** The number of routines **to be executed** in any time interval $[t, t+x]$ is identical to that in $[t+h, t+h+x]$

Priority-Driven Scheduling

System Model

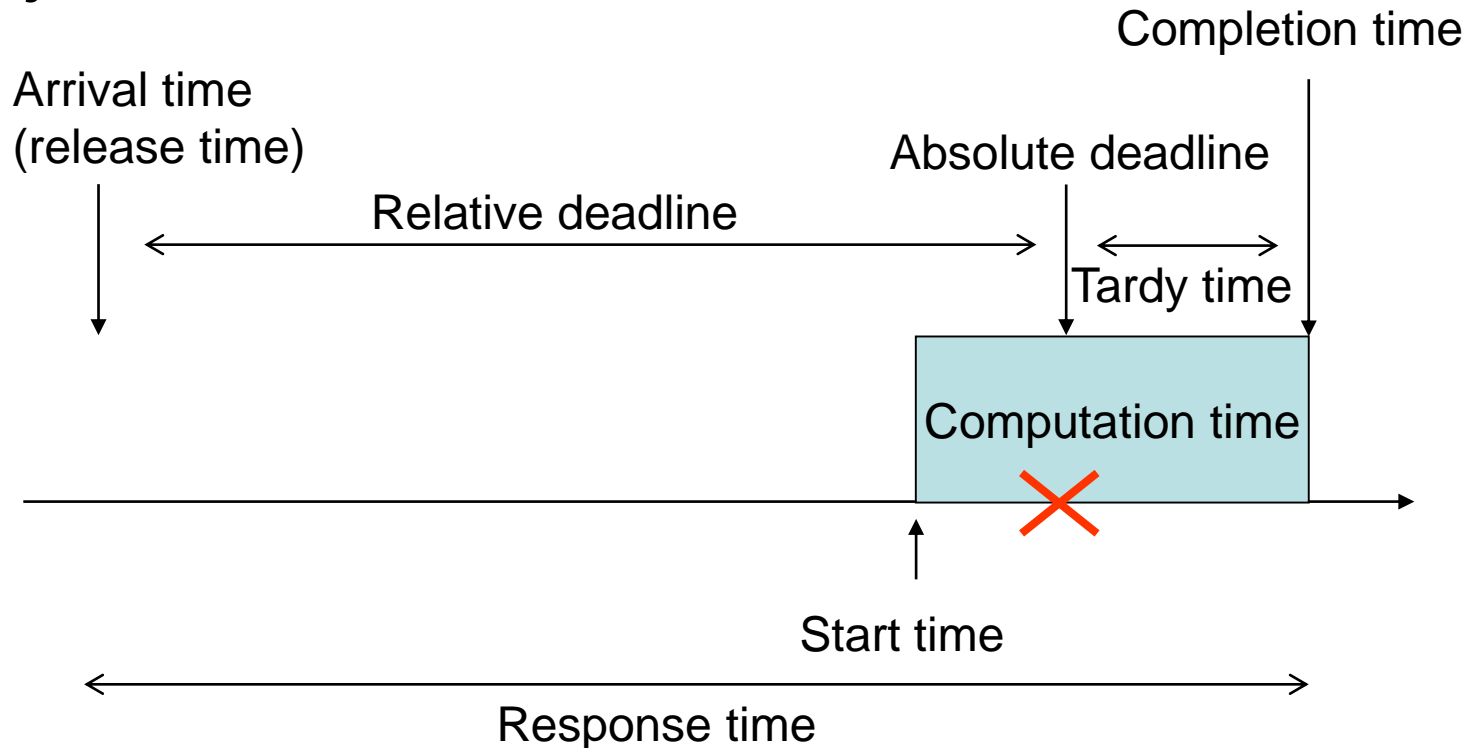
- A job with real-time constraints



The job completes before its deadline, that means the deadline is satisfied.

System Model

- A job with real-time constraints




The job completes after its deadline, that means the deadline is **violated** or an **overflow** occurs.

System Model

- A task set is of a number of tasks
 - $\{T_1, T_2, \dots, T_n\}$
 - Tasks share nothing but CPU, and tasks are independent to one another
- A task T_i is a template of jobs, where jobs refer to recurrences of task.
 - Every job executes the same piece of code
 - Of course, different input and run-time conditions cause jobs behaves differently
 - $J_{i,j}$ refers to the j -th job of task T_i
 - The computation time c_i of jobs is bounded and known a priori

System Model

- A purely periodic task
 - Jobs of a task **T** recur every fixed time interval **p**
 - A job must be completed before the next job arrives
 - Relative deadlines for jobs are, implicitly, the period
 - T is defined as (c,p)



Periodic task $T_1 = (3, 6)$

System Model


- Priority
 - Reflect the urgency of jobs
 - Any job inherits its task's priority
- Preemptivity
 - As a high-priority task arrives, it preempts the execution of any low-priority tasks

System Model

- Checklist
 - Periodic tasks
 - Real-time constraints
 - Priority
 - Preemptivity

Rate-Monotonic Scheduling

- Task-level fixed-priority scheduling
 - All jobs inherit its task's priority
 - Usually abbreviated as fixed-priority scheduling
- Tasks' priorities are inversely proportional to their period lengths




Rate-Monotonic Scheduling

- Critical instant (critical instance) of task T_i
 - A job $J_{i,c}$ of task T_i released at T_i 's critical instant would have the longest response time
 - $J_{i,c}$ would be the one that is “hardest” to meet its deadline
 - If $J_{i,c}$ succeeds in satisfying its deadline, then any job of T_i always succeeds for any cases
 - Since in any other cases deadlines are easier to meet

Rate-Monotonic Scheduling

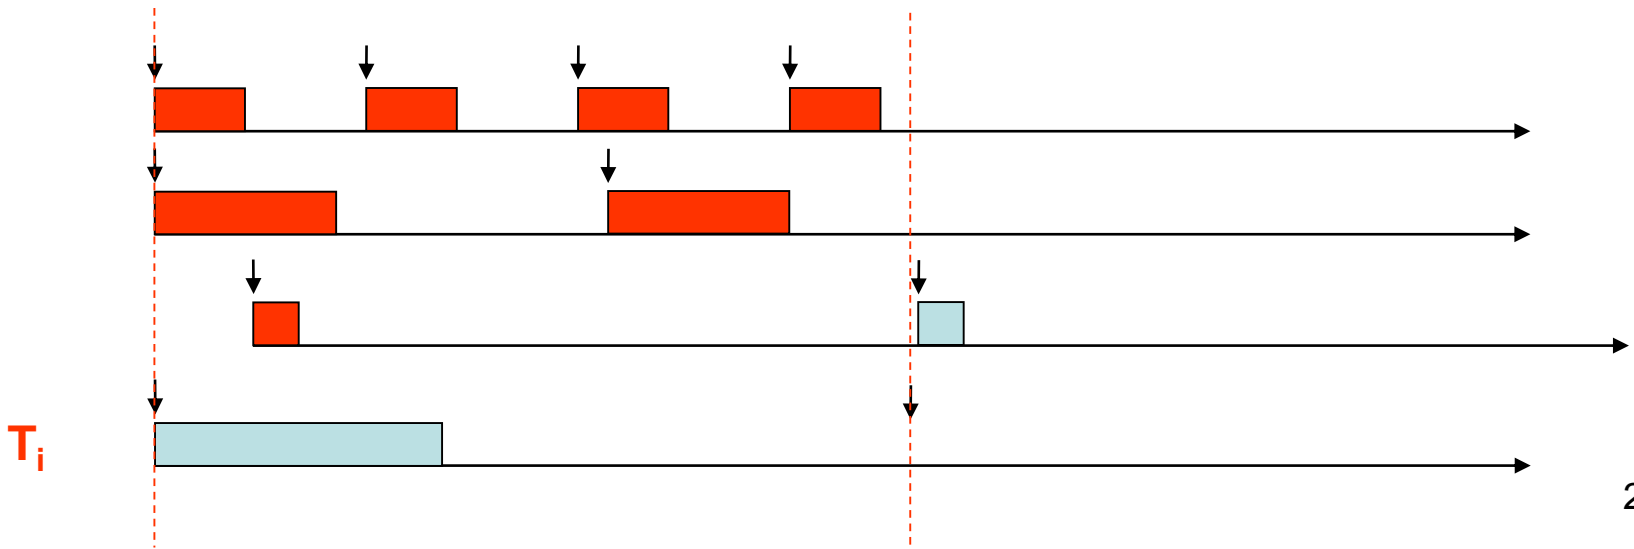
- **Theorem:** A critical instant of any task T_i occurs when one of its job $J_{i,c}$ is released at the same time with a job of every higher-priority task (i.e., in-phase).



Rate-Monotonic Scheduling


- Proof: “interferences” from high-priority tasks is the maximum within the first period of T_i

$$\sum_{j < i} c_j \left\lceil \frac{p_i}{p_j} \right\rceil$$



Rate-Monotonic Scheduling

- Critical instant: Why ceiling function?




Rate-Monotonic Scheduling

- Response time analysis
 - The response time of the job of T_i at critical instant can be calculated by the following recursive function

$$r_0 = \sum_{\forall i} c_i$$

$$r_n = \sum_{\forall i} c_i \left\lceil \frac{r_{n-1}}{p_i} \right\rceil$$



- Observation: the sequence of r_x , $x \geq 0$ may or may not converge

Rate-Monotonic Scheduling

- **Theorem:** Given a task set $\{T_1, T_2, \dots, T_n\}$, if at critical instant the response time of the first job of task T_i , for each i , converges no later than p_i , then jobs never miss their deadlines
- Observations
 - If the task set survives critical instant, then it will survive any task phasing
 - The analysis is an exact schedulability test for RMS
 - Usually referred to as “Rate-Monotonic Analysis”, RMA for short

Rate-Monotonic Scheduling

- Example: $T1=(2,5)$, $T2=(2,7)$, $T3=(3,8)$
 - T1:
 - $R_0=2 \leq 5$ ok
 - T2:
 - $R_0=2+2=4 \leq 7$
 - $R_1=2 * \lceil 4/5 \rceil + 2 * \lceil 4/7 \rceil = 4 \leq 7$ ok
 - T3:
 - $R_0=2+2+3=7 \leq 8$
 - $R_1= 2 * \lceil 7/5 \rceil + 2 * \lceil 7/7 \rceil + 3 * \lceil 7/8 \rceil = 9 > 8$ failed
 - Note: each task succeeds \rightarrow the task set

Rate-Monotonic Scheduling

- Proof:
 - If the response time converges at r_n , then the first lowest-priority job completes at r_n
 - If r_n is before p_n , then the first lowest-priority job meets its deadline if critical instant occurs
 - Since the job survives critical instant, it always succeed satisfying its deadline under any task phasing

Rate-Monotonic Scheduling

- Test every T_i for schedulability!!
 - $\{T1=(3,6), T2=(3.1,9), T3=(1,18)\}$
 - Response analysis of $T3$:
 - $R0=7.1, R1=10.1, R2=13.2, R3=16.2, R4=16.2 < 18$
 - $\{T1, T2, T3\}$ is schedulable!?
 - However, $\{T1, T2\}$ is not schedulable!!!

Rate-Monotonic Scheduling

- Computational complexity
 - $O(n^2 * p_n)$, pseudo-polynomial time
 - Would be *extremely slow* when periods of tasks are small and prime to one another
 - Would be very fast when periods are harmonically related

Rate-Monotonic Scheduling

- Phenomena
 - Even though RMA is an exact test for fixed-priority scheduling, it is not often used, especially for those dynamic systems
 - RMA is more suitable for static systems
 - Are there any schedulability tests being efficient enough for on-line implementation?
 - No slower than polynomial time

Rate-Monotonic Scheduling

- A trivial schedulability test
 - The system accepts a task set T if the following conditions are both true
 - There are no other tasks in the system
 - $c_1/p_1 \leq 1$
 - The algorithm is efficient enough (i.e., $O(1)$)
 - Too conservative!! Very Poor CPU utilization!!

Rate-Monotonic Scheduling

- Definition

- Utilization factor of task $T=(c,p)$ is defined as

$$\frac{c}{p}$$

- CPU utilization of a task set $\{T_1, T_2, \dots, T_n\}$ is

$$U = \sum_{i=1}^n \frac{c_i}{p_i}$$

- Observation: if the total utilization exceeds 1 then the task set is not schedulable

Rate-Monotonic Scheduling

- **Theorem:** [LL73] Given a task set $\{T_1, T_2, \dots, T_n\}$. It is schedulable by RMS if

$$\sum_{i=1}^n \frac{c_i}{p_i} \leq U(n) = n(2^{1/n} - 1)$$

- Observation:
 - If the test succeeds then the task is schedulable
 - A sufficient condition for schedulability

Rate-Monotonic Scheduling

- Proof: Let us consider two tasks only



$$C_1 \leq P_2 - P_1(\lfloor P_2/P_1 \rfloor)$$

The largest possible C2 is

$$P_2 - C_1(\lceil P_2/P_1 \rceil)$$

Total utilization factor is


$$U = 1 + \underline{C_1(1/P_1 - (1/P_2)(\lceil P_2/P_1 \rceil))}$$

T2's 2nd job does not overlap the immediately preceding job of T1

- U monotonically decreases with C_1
 - Because U never greater than 1, so the rightmost term in the last equation is always negative

Rate-Monotonic Scheduling

- Proof: cont'd.



$$C_1 \geq P_2 - P_1(\lfloor P_2/P_1 \rfloor)$$

The largest possible C_2 is

$$\dots - C_1(\lfloor P_2/P_1 \rfloor) + P_1(\lfloor P_2/P_1 \rfloor)$$

Total utilization factor is

$$U = (P_1/P_2)\lfloor P_2/P_1 \rfloor + C_1((1/P_1) - (1/P_2)(\lfloor P_2/P_1 \rfloor))$$

T2's 2nd job overlaps the immediately preceding job of T1


- U monotonically increases with C_1

Rate-Monotonic Scheduling

- Proof: Cont'd.

- It can be found that the minimal U occurs at

$$C_1 = P_2 - P_1 \lfloor P_2 / P_1 \rfloor$$




- By some differentiation, the minimal achievable utilization is

$$U(2) = 2(2^{1/2} - 1)$$

Rate-Monotonic Scheduling

- To find “the smallest” among “the largest achievable processor utilizations that can be achieved by different task sets”




Rate-Monotonic Scheduling

- Simon says: To generalize the proof to n tasks is easy :)
- If a task set of n tasks has a total utilization being no greater than $U(n)$, then it is guaranteed to be schedulable by RMS
 - Because the most hard-to-schedule task set having the same total utilization is schedulable
 - The test's time complexity is $O(n)$, which is very efficient for on-line implementation

Rate-Monotonic Scheduling


- When $x \rightarrow$ infinitely large, $U(x) \rightarrow 0.68$




1	1
2	0.828427
3	0.779763
4	0.756828
5	0.743492
6	0.734772
7	0.728627
8	0.724062
9	0.720538
10	0.717735
11	0.715452
12	0.713557
13	0.711959
14	0.710593
15	0.709412

Rate-Monotonic Scheduling

- Example 1: (1,3), (2,5)
 - Utilization = $0.73 \leq U(2) = 0.828$




- Example 2: (4.5,9), (3,6)
 - Utilization = $100\% > U(2) = 0.828$



Rate-Monotonic Scheduling


- Example 3: $(1,2), (2,4)$
 - Utilization = 100% > $U(2) = 0.828$



- Example 2 and 3 shows that, we know nothing about those task sets of total utilization > the utilization bound!
- But we do know those \leq the utilization bound is schedulable!

Rate-Monotonic Scheduling

- Sufficiency but no necessity
 - Utilization test provides a fast way to check if a task set is schedulable
 - Any task set fails utilization test, does not implies that it is not schedulable



U : universe of task sets

$\sim S$: task sets unschedulable by RMS

- Example 2

S : task sets schedulable by RMS

- Example 1 and Example 3

A : Those can be found by utilization test

- Example 1

Example 3 is in $S-A$

Rate-Monotonic Scheduling

- Summary
 - Explicit prioritization over tasks
 - To **decide** task sets' schedulability is costly
 - Sufficient tests were developed for fast admission control

Priority-Driven Scheduling: Dynamic-Priority Scheduling

Earliest-Deadline-First Scheduling

- Definition
 - Feasible
 - A set of tasks is said to be feasible if there is some way to schedule the tasks without any deadline violations
 - Schedulable
 - Given a scheduling algorithm A
 - A set of tasks is said to be schedulable if algorithm A successfully schedule the tasks without any deadline violations
- Observations
 - A feasible task set may not be schedulable by RMS
 - If a task set is schedulable by some algorithm A, then it is feasible

Earliest-Deadline-First Scheduling

- EDF scheduling algorithm always pick a pending job which has the earliest deadline for execution
 - A job having an earlier deadline is assigned to a higher “priority”
 - Priority in EDF is not a task-wide notion
 - Jobs of a task may have different “priorities”
 - But due to the relative deadline of a job never changes, EDF can be classified as a job-level fixed-priority scheduling
 - You’d better to avoid using the term “priority” for EDF since there is no explicit definition


Earliest-Deadline-First Scheduling

- If an algorithm schedulable \leftrightarrow feasible
 - It can be referred to as a universal scheduling algorithm
- What is the **universal scheduling algorithm** for periodic and preemptive uniprocessor systems?
 - EDF!


Earliest-Deadline-First Scheduling

- Example

Not scheduable by RMS




Schedulable by EDF



Earliest-Deadline-First Scheduling

- ***Theorem:*** With EDF, there is no idle time before an overflow
- **Observation:** A very strong statement that implies optimality of EDF in terms of schedulability

Proof




- Suppose that there is an overflow at time t_3 , and the processor idles between t_1 and t_2
- If we move "a" forward to be aligned with t_2 , the overflow would occur earlier than it was (i.e., at or before t_3)
 - That is because EDF's discipline: moving forward means promoting the urgency of T_1 's jobs
- By repeating the above action, jobs a, b, and c can be aligned at t_2
 - → that contradicts the assumption! From t_2 on, there is no idle until the overflow

Earliest-Deadline-First Scheduling

- ***Theorem:*** A set of tasks is schedulable by EDF if and only if its total CPU utilization is no more than 1
- Observation: \rightarrow is easy, \leftarrow requires some reasoning similar to the proof of the last theorem

overflow



→: suppose that $U \leq 1$ but the system is not schedulable by EDF

- Suppose that there is an overflow at time T
 - Jobs a 's have deadlines at time T
 - Job b 's have deadlines after time T

• Case A: none of job b 's is executed before T

- The total computational demand between $[0, T]$ is

$$C_1(\lfloor T/P_1 \rfloor) + C_2(\lfloor T/P_2 \rfloor) + \dots + C_n(\lfloor T/P_n \rfloor)$$

- Since there is no idle before an overflow


$$C_1(\lfloor T/P_1 \rfloor) + C_2(\lfloor T/P_2 \rfloor) + \dots + C_n(\lfloor T/P_n \rfloor) > T$$

- That implies $U > 1$

• → ←

Boxes: “arrival” times!!

overflow



Suppose b2 is
being executed
here...

Case B: some of job b's are executed before T

- Because an overflow occurs at T, the violated jobs must be a's
 - Right before T, there must be some job a's being executed
 - Let in $[T', T]$ there is no job b's being executed
- Just before T', some of b's is being executed! (the definition of T')
 - It means that all jobs have deadlines $\leq T$ and arrive before T' have completed before T'
- Back to $[T', T]$, the total computation demand is no less than

$$C_1(\lfloor T - T'/P_1 \rfloor) + C_2(\lfloor T - T'/P_2 \rfloor) + \dots + C_n(\lfloor T - T'/P_n \rfloor)$$
 - Because there is deadline violations, so

$$C_1(\lfloor T - T'/P_1 \rfloor) + C_2(\lfloor T - T'/P_2 \rfloor) + \dots + C_n(\lfloor T - T'/P_n \rfloor) \geq T - T'$$

Earliest-Deadline-First Scheduling

- Summary
 - A universal scheduling algorithm for real-time periodic tasks
 - Urgency of tasks is dynamic
 - But static for jobs
 - Job-level fixed-priority scheduling

Comparison

	RMS	EDF
Optimality	Optimal for fixed-priority scheduling	Universal
Schedulability test	Exact test is slow (PP), conservative tests are adopted	$O(n)$ for exact test
Algorithm time complexity	$O(1)$ job insertion is possible	Both job insertion and dispatch take $O(\log n)$ time
Overload survivability	High and predictable	Low and unmanageable**
Responsiveness	High priority tasks always have shorter response time	Non-intuitive to reach conclusions
Ease of implementation	Pretty simple	Relatively complicated
Run-time overheads (like preemption)	Low	High

Is it true?